

Pgloader Overview and Features

what is Pgloader ? how it works internally ? what are the characteristics of pgloader ? what are the features it has ?

explain all about pgloader so that i donot have any doubt in this and have cristal and clear knowledge of pgloader

Below is a complete, crystal-clear, zero-doubt explanation of pgloader, written in simple but deep technical language, step-by-step.

I'll explain what it is → why it exists → how it works internally → characteristics → features → real-world behavior → when to use & when not.

1 What is pgloader?

pgloader is an open-source data migration tool used to load data into PostgreSQL from other databases, files, or formats.

👉 In simple words:

pgloader helps you move data from one system into PostgreSQL safely, fast, and automatically.

Common use cases

- MySQL → PostgreSQL
- MSSQL → PostgreSQL
- SQLite → PostgreSQL
- CSV / TSV files → PostgreSQL
- Old legacy DB → PostgreSQL

2 Why pgloader exists (The core problem it solves)

Migrating databases is hard, because:

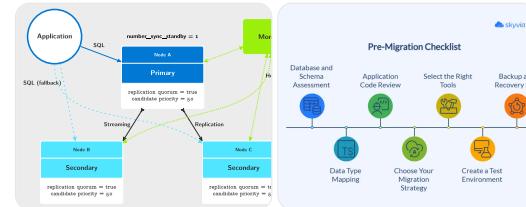
- Different databases have different data types
- Constraints behave differently

- Indexes and foreign keys cause failures
- Large datasets are slow to migrate
- Manual scripts break easily

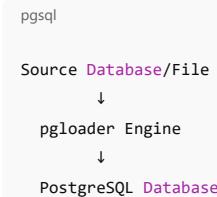
pgloader solves this by:

- ✓ Automatically converting data types
- ✓ Loading data in parallel (very fast)
- ✓ Handling constraints smartly
- ✓ Retrying, logging, and validating

3 High-level working flow of pgloader



Simple flow



4 How pgloader works internally (DEEP EXPLANATION)

This is the most important part 👉

I'll explain it like inside-the-engine behavior.

◆ Step 1: Source Connection

pgloader first connects to the **source system**:

- MySQL / MSSQL / SQLite
- Or file system (CSV, TSV)

It reads:

- Tables
- Columns
- Data types
- Constraints
- Index definitions

👉 This metadata is stored internally in memory.

◆ Step 2: Schema Analysis & Type Mapping

pgloader does not blindly copy data.

Instead, it:

- Maps source data types → PostgreSQL types

Example:

Source DB	Source Type	PostgreSQL Type
MySQL	INT	INTEGER
MySQL	TINYINT(1)	BOOLEAN
MySQL	DATETIME	TIMESTAMP
MSSQL	NVARCHAR	TEXT

This mapping is:

- ✓ Automatic
- ✓ Configurable (you can override it)

◆ Step 3: Target Preparation (VERY IMPORTANT)

Before inserting data, pgloader prepares PostgreSQL:

- ✓ Creates tables
- ✓ Creates sequences
- ✓ Temporarily disables constraints

- Foreign keys
- Triggers
- Indexes

👉 Why?

Because inserting millions of rows with constraints ON is slow and error-prone.

◆ Step 4: Parallel Data Loading (Key Performance Feature)

pgloader loads data using:

- Multiple worker threads
- COPY command (fastest PostgreSQL insertion)

Internally:

mathematica

```
Thread 1 → Table A
Thread 2 → Table B
Thread 3 → Table C
```

👉 This is why pgloader is much faster than manual INSERT scripts.

◆ Step 5: Data Transformation (On-the-fly)

During loading, pgloader can:

- Convert encodings
- Fix invalid values
- Apply expressions
- Rename columns
- Skip bad rows (optional)

Example:

lisp

```
CAST type datetime TO timestampz
```

◆ Step 6: Post-Load Cleanup

After data is loaded:

- ✓ Re-enable constraints
- ✓ Create indexes
- ✓ Validate foreign keys
- ✓ Analyze tables (`ANALYZE`)

Now PostgreSQL is ready for production use.

5 Characteristics of pgloader (Core Nature)

These define how pgloader behaves.

◆ 1. PostgreSQL-Focused

- Built specifically for PostgreSQL
- Uses PostgreSQL internals like:
 - `COPY`
 - `ANALYZE`
 - Sequences

◆ 2. Declarative Configuration

You describe what you want, not how.

Example:

```
sql
LOAD DATABASE
  FROM mysql://user:pass@localhost/db
  INTO postgresql://user:pass@localhost/db;
```

◆ 3. Fault-Tolerant

- Can skip errors
- Logs rejected rows
- Supports retries

◆ 4. High Performance

- Parallel loading
- Bulk `COPY`
- Constraint optimization

◆ 5. Non-Destructive

- Does NOT modify source DB
- Only reads from it

6 Features of pgloader (What it can do)

✓ Database to PostgreSQL

- MySQL → PostgreSQL
- MSSQL → PostgreSQL
- SQLite → PostgreSQL

✓ File-based Loading

- CSV
- TSV
- Fixed-width files

✓ Automatic Schema Creation

- Tables
- Sequences
- Indexes

Data Type Conversion

- Automatic mapping
- Custom rules

Constraint Handling

- Disable → load → enable
- Avoids FK failures

Parallelism

- Multiple tables at once
- Massive speed boost

Encoding Handling

- UTF-8 fixes
- Character set conversion

Logging & Monitoring

- Progress output
- Error reports
- Rejected rows file

Scriptable & Repeatable

- Same script can be run again
- Ideal for CI/CD migrations

7 Real-World Example (Conceptual)

MySQL → PostgreSQL migration

pgsql

1 million users
50 tables
Foreign keys everywhere

Without pgloader

- ✗ Manual scripts
- ✗ FK errors
- ✗ Takes hours or days

With pgloader

- ✓ 1 command
- ✓ Handles FK automatically
- ✓ Finishes in minutes

8 When should you use pgloader?

Use pgloader when:

- ✓ Migrating to PostgreSQL
- ✓ Large datasets
- ✓ Schema + data both needed
- ✓ Speed matters

Do NOT use pgloader when:

- ✗ PostgreSQL is NOT the target
- ✗ You need real-time replication
- ✗ You need bi-directional sync

(pgloader is for **migration**, not replication)

9 pgloader vs Liquibase (Important clarity)

pgloader	Liquibase
Data migration	Schema versioning
One-time bulk load	Continuous DB evolution
External sources	PostgreSQL only
High-volume data	Small incremental changes

👉 They solve different problems.

10 Mental Model (Remember this)

pgloader is a migration engine that reads, transforms, and bulk-loads data into PostgreSQL with speed and safety.

If you understand:

- ✓ Source reading
- ✓ Type mapping
- ✓ Parallel COPY
- ✓ Constraint management

👉 You fully understand pgloader.

If you want, next I can:

- Explain pgloader vs AWS DMS
- Show real pgloader scripts
- Explain error handling scenarios
- Explain pgloader in production migration

Just tell me 👍

how thread works in pgloader

Below is a **deep, crystal-clear explanation of how threads work in pgloader**, written so that you won't have any confusion, even internally/architecturally.

I'll go step-by-step and also give you a **mental model** you can always remember.

1 What does "thread" mean in pgloader?

In pgloader, a **thread** is an independent worker that performs part of the migration in parallel.

👉 Each thread:

- Reads data from the source
- Transforms it (if needed)
- Loads it into PostgreSQL using `COPY`

So instead of:

css

```
Table A → done
Table B → done
Table C → done
```

pgloader does:

mathematica

```
Table A → Thread 1
Table B → Thread 2
Table C → Thread 3
```

This is the **core reason pgloader is fast**.

2 Why pgloader needs threads (the problem it solves)

Without threads:

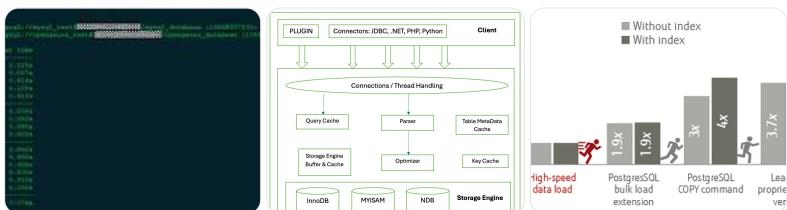
- One table loads at a time
- CPU is mostly idle
- Network is underused

- Migration is slow

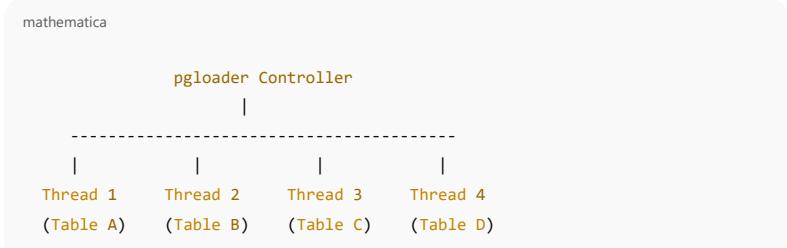
With threads:

- ✓ CPU fully utilized
- ✓ Network fully utilized
- ✓ Disk IO parallelized
- ✓ PostgreSQL COPY speed maximized

3 High-level threading architecture



Think of pgloader like this:



Each thread is **independent** but **coordinated**.

4 Internal thread lifecycle (VERY IMPORTANT)

Let's follow one pgloader thread from birth to death.

◆ Step 1: Thread creation

pgloader:

- Detects how many tables need to be loaded
- Creates a **thread pool**

Default:

ini

`workers = number of CPU cores`

(You can override this.)

◆ Step 2: Task assignment

Each thread is assigned a **task**, such as:

- Load one table
- Load part of a large table (chunking)

Example:

mathematica

Thread 1 → users
Thread 2 → orders
Thread 3 → products

◆ Step 3: Source reading (per thread)

Each thread:

- Opens its **own source connection**
- Reads rows in batches
- Streams data (does NOT load entire table in memory)

👉 This avoids memory explosion.

◆ Step 4: Data transformation (inside thread)

Inside each thread:

- Data types are converted

- NULLs fixed
- Encoding corrected
- Invalid values handled

This happens **row-by-row, inside the thread**.

◆ Step 5: COPY into PostgreSQL

This is the most critical part.

Each thread:

- Opens its own PostgreSQL connection
- Executes:

sql

```
COPY table_name FROM STDIN;
```

So internally:

powershell

```
Thread 1 → COPY users
Thread 2 → COPY orders
Thread 3 → COPY products
```

👉 PostgreSQL handles multiple COPY operations **concurrently**.

◆ Step 6: Thread completion

After finishing:

- Thread reports success/failure
- Releases DB connections
- Returns to thread pool (or exits)

5 Two levels of parallelism in pgloader

This is where most people get confused.

◆ Level 1: Table-level parallelism

Different tables loaded at the same time.

mathematica

```
users → Thread 1
orders → Thread 2
items → Thread 3
```

◆ Level 2: Chunk-level parallelism (large tables)

If a table is **huge**, pgloader can:

- Split it into chunks
- Load chunks in parallel

Example:

mathematica

```
users (10M rows)
└─ chunk 1 → Thread 1
└─ chunk 2 → Thread 2
└─ chunk 3 → Thread 3
```

6 How pgloader avoids thread conflicts

Parallelism is dangerous if unmanaged. pgloader avoids problems by:

Constraint handling

- Foreign keys disabled before loading
- Re-enabled after all threads finish

Index handling

- Indexes created **after** data load
- No index contention during COPY

Transaction isolation

- Each thread uses its own transaction

- No shared locks on source

7 Thread configuration in pgloader

You can control threads explicitly.

Example:

```
lisp
WITH
  workers = 4,
  concurrency = 2
```

Meaning:

- `workers` → total worker threads
- `concurrency` → how many tables at the same time

8 What happens if threads are too many?

✗ Too many threads causes:

- PostgreSQL connection exhaustion
- Disk IO saturation
- Slower migration (counter-productive)

✓ Best practice:

```
nginx
threads = CPU cores
```

9 Failure handling with threads

If one thread fails:

- Other threads continue
- Error is logged
- Migration does NOT crash completely (unless fatal)

This is called **partial fault tolerance**.

10 Mental model (remember forever)

pgloader threads are independent workers that read, transform, and COPY data in parallel, while a central controller coordinates constraints and order.

If you understand:

- ✓ Thread pool
- ✓ One thread = one COPY stream
- ✓ Parallel tables / chunks
- ✓ Constraints disabled globally

👉 You have **complete mastery of pgloader threading**.

11 Short real-world analogy

Imagine:

- 1 warehouse (PostgreSQL)
- 5 trucks (threads)
- Each truck brings goods (data)
- Security gates (constraints) are opened after all trucks arrive

That's exactly pgloader threading.

If you want next, I can:

- Explain **workers vs concurrency** with diagrams
- Explain **pgloader performance tuning**
- Explain **thread behavior during FK-heavy schemas**
- Compare **pgloader threads vs Java threads**

Just tell me 👍