Pgloader Overview and Features

what is Pgloader ? how it works internally ? what are the characterstics of pgloader ? what are the features it has ?

explain all about pgloader so that i donot have any doubt in this and have cristal and clear knowledge of pgloader

Below is a **complete, crystal-clear, zero-doubt explanation of pgloader**, written in **simple but deep technical language**, step-by-step.
I'll explain **what it is → why it exists → how it works internally → characteristics → features → real-world behavior → when to use & when not**.

1️⃣ What is pgloader?

pgloader is an **open-source data migration tool** used to **load data into PostgreSQL** from **other databases, files, or formats**.

👉 In simple words:

> **pgloader helps you move data from one system into PostgreSQL safely, fast, and automatically.**

Common use cases

- MySQL → PostgreSQL
- MSSQL → PostgreSQL
- SQLite → PostgreSQL
- CSV / TSV files → PostgreSQL
- Old legacy DB → PostgreSQL

2️⃣ Why pgloader exists (The core problem it solves)

Migrating databases is **hard**, because:

- Different databases have **different data types**
- Constraints behave differently

- Indexes and foreign keys cause failures
- Large datasets are **slow** to migrate
- Manual scripts break easily

pgloader solves this by:

- ✔ Automatically converting data types
- ✔ Loading data **in parallel** (very fast)
- ✔ Handling constraints smartly
- ✔ Retrying, logging, and validating

3️⃣ High-level working flow of pgloader

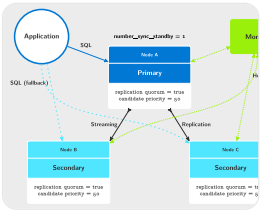

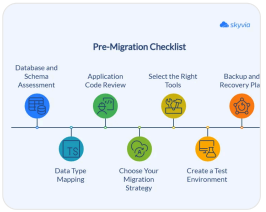

Simple flow

```pgsql
Source Database/File
        ↓
  pgloader Engine
        ↓
  PostgreSQL Database
```

4️⃣ How pgloader works internally (DEEP EXPLANATION)

This is the most important part 👇
I'll explain it like **inside-the-engine behavior**.

🔹 Step 1: Source Connection

pgloader first connects to the **source system**:

- MySQL / MSSQL / SQLite
- Or file system (CSV, TSV)

It reads:

- Tables
- Columns
- Data types
- Constraints
- Index definitions

👉 This metadata is stored **internally in memory**.

🔹 Step 2: Schema Analysis & Type Mapping

pgloader **does not blindly copy data**.

Instead, it:

- Maps source data types → PostgreSQL types

Example:

Source DB	Source Type	PostgreSQL Type
MySQL	`INT`	`INTEGER`
MySQL	`TINYINT(1)`	`BOOLEAN`
MySQL	`DATETIME`	`TIMESTAMP`
MSSQL	`NVARCHAR`	`TEXT`

This mapping is:
✔ Automatic
✔ Configurable (you can override it)

🔹 Step 3: Target Preparation (VERY IMPORTANT)

Before inserting data, pgloader prepares PostgreSQL:

✔ Creates tables
✔ Creates sequences
✔ Temporarily **disables constraints**

- Foreign keys
- Triggers
- Indexes

👉 Why?
Because inserting millions of rows with constraints ON is **slow and error-prone**.

🔹 Step 4: Parallel Data Loading (Key Performance Feature)

pgloader loads data using:

- **Multiple worker threads**
- **COPY command** (fastest PostgreSQL insertion)

Internally:

```mathematica
Thread 1 → Table A
Thread 2 → Table B
Thread 3 → Table C
```

👉 This is why pgloader is **much faster than manual INSERT scripts**.

🔹 Step 5: Data Transformation (On-the-fly)

During loading, pgloader can:

- Convert encodings
- Fix invalid values
- Apply expressions
- Rename columns
- Skip bad rows (optional)

Example:

```lisp
```

```
CAST type datetime TO timestamptz
```

🔹 Step 6: Post-Load Cleanup

After data is loaded:

- ✔ Re-enable constraints
- ✔ Create indexes
- ✔ Validate foreign keys
- ✔ Analyze tables (`ANALYZE`)

Now PostgreSQL is **ready for production use**.

5️⃣ Characteristics of pgloader (Core Nature)

These define **how pgloader behaves**.

🔸 1. PostgreSQL-Focused

- Built **specifically** for PostgreSQL
- Uses PostgreSQL internals like:
  - COPY
  - ANALYZE
  - Sequences

🔸 2. Declarative Configuration

You describe **what you want**, not how.

Example:

```sql
LOAD DATABASE
     FROM mysql://user:pass@localhost/db
     INTO postgresql://user:pass@localhost/db;
```

🔸 3. Fault-Tolerant

- Can skip errors
- Logs rejected rows
- Supports retries

🔸 4. High Performance

- Parallel loading
- Bulk COPY
- Constraint optimization

🔸 5. Non-Destructive

- Does NOT modify source DB
- Only reads from it

6️⃣ Features of pgloader (What it can do)

✅ Database to PostgreSQL

- MySQL → PostgreSQL
- MSSQL → PostgreSQL
- SQLite → PostgreSQL

✅ File-based Loading

- CSV
- TSV
- Fixed-width files

✅ Automatic Schema Creation

- Tables
- Sequences
- Indexes

✅ Data Type Conversion

- Automatic mapping
- Custom rules

✅ Constraint Handling

- Disable → load → enable
- Avoids FK failures

✅ Parallelism

- Multiple tables at once
- Massive speed boost

✅ Encoding Handling

- UTF-8 fixes
- Character set conversion

✅ Logging & Monitoring

- Progress output
- Error reports
- Rejected rows file

✅ Scriptable & Repeatable

- Same script can be run again
- Ideal for CI/CD migrations

7️⃣ Real-World Example (Conceptual)

MySQL → PostgreSQL migration


```pgsql
1 million users
50 tables
Foreign keys everywhere
```

Without pgloader

❌ Manual scripts
❌ FK errors
❌ Takes hours or days

With pgloader

✔ 1 command
✔ Handles FK automatically
✔ Finishes in minutes

8️⃣ When should you use pgloader?

✅ Use pgloader when:

✔ Migrating **to PostgreSQL**
✔ Large datasets
✔ Schema + data both needed
✔ Speed matters

❌ Do NOT use pgloader when:

❌ PostgreSQL is NOT the target
❌ You need real-time replication
❌ You need bi-directional sync

(pgloader is for **migration**, not replication)

9️⃣ pgloader vs Liquibase (Important clarity)

pgloader	Liquibase
Data migration	Schema versioning
One-time bulk load	Continuous DB evolution
External sources	PostgreSQL only
High-volume data	Small incremental changes

👉 They **solve different problems**.

🔟 Mental Model (Remember this)

> **pgloader is a migration engine that reads, transforms, and bulk-loads data into PostgreSQL with speed and safety.**

If you understand:
✔ Source reading
✔ Type mapping
✔ Parallel COPY
✔ Constraint management

👉 You **fully understand pgloader**.

If you want, next I can:

- Explain **pgloader vs AWS DMS**
- Show **real pgloader scripts**
- Explain **error handling scenarios**
- Explain **pgloader in production migration**

Just tell me 👍

how thread works in pgloader

Below is a **deep, crystal-clear explanation of how threads work in pgloader**, written so that **you won't have any confusion**, even internally/architecturally.

I'll go step-by-step and also give you a **mental model** you can always remember.

1️⃣ What does "thread" mean in pgloader?

In **pgloader**, a **thread** is an **independent worker** that performs part of the migration **in parallel**.

👉 Each thread:

- Reads data from the source
- Transforms it (if needed)
- Loads it into PostgreSQL using `COPY`

So instead of:

```
css

Table A → done
Table B → done
Table C → done
```

pgloader does:

```
mathematica

Table A → Thread 1
Table B → Thread 2
Table C → Thread 3
```

This is the **core reason pgloader is fast**.

2️⃣ Why pgloader needs threads (the problem it solves)

Without threads:

- One table loads at a time
- CPU is mostly idle
- Network is underused

- Migration is slow

With threads:

- ✔ CPU fully utilized
- ✔ Network fully utilized
- ✔ Disk IO parallelized
- ✔ PostgreSQL COPY speed maximized

3️⃣ High-level threading architecture

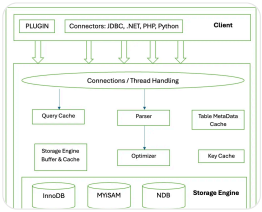

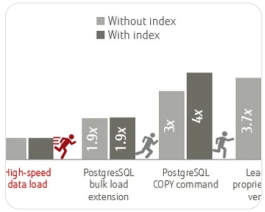

Think of pgloader like this:

```
mathematica

                pgloader Controller
                        |
    ----------------------------------------
    |           |           |           |
Thread 1    Thread 2    Thread 3    Thread 4
(Table A)   (Table B)   (Table C)   (Table D)
```

Each thread is **independent** but **coordinated**.

4️⃣ Internal thread lifecycle (VERY IMPORTANT)

Let's follow **one pgloader thread from birth to death**.

🔹 Step 1: Thread creation

pgloader:

- Detects how many tables need to be loaded
- Creates a **thread pool**

Default:

```
ini

workers = number of CPU cores
```

(You can override this.)

🔹 Step 2: Task assignment

Each thread is assigned a **task**, such as:

- Load one table
- Load part of a large table (chunking)

Example:

```
mathematica

Thread 1 → users
Thread 2 → orders
Thread 3 → products
```

🔹 Step 3: Source reading (per thread)

Each thread:

- Opens its **own source connection**
- Reads rows in batches
- Streams data (does NOT load entire table in memory)

👉 This avoids memory explosion.

🔹 Step 4: Data transformation (inside thread)

Inside each thread:

- Data types are converted

- NULLs fixed
- Encoding corrected
- Invalid values handled

This happens **row-by-row**, **inside the thread**.

🔹 Step 5: COPY into PostgreSQL

This is the most critical part.

Each thread:

- Opens its **own PostgreSQL connection**
- Executes:

```sql
COPY table_name FROM STDIN;
```

So internally:

```powershell
Thread 1 → COPY users
Thread 2 → COPY orders
Thread 3 → COPY products
```

👉 PostgreSQL handles multiple COPY operations **concurrently**.

🔹 Step 6: Thread completion

After finishing:

- Thread reports success/failure
- Releases DB connections
- Returns to thread pool (or exits)

5️⃣ Two levels of parallelism in pgloader

This is where most people get confused.

🔸 Level 1: Table-level parallelism

Different tables loaded at the same time.

```mathematica
users   → Thread 1
orders  → Thread 2
items   → Thread 3
```

🔸 Level 2: Chunk-level parallelism (large tables)

If a table is **huge**, pgloader can:

- Split it into chunks
- Load chunks in parallel

Example:

```mathematica
users (10M rows)
  ├─ chunk 1 → Thread 1
  ├─ chunk 2 → Thread 2
  ├─ chunk 3 → Thread 3
```

6️⃣ How pgloader avoids thread conflicts

Parallelism is dangerous if unmanaged. pgloader avoids problems by:

✅ Constraint handling

- Foreign keys disabled before loading
- Re-enabled after all threads finish

✅ Index handling

- Indexes created **after** data load
- No index contention during COPY

✅ Transaction isolation

- Each thread uses its own transaction

- No shared locks on source

7️⃣ Thread configuration in pgloader

You can control threads explicitly.

Example:

```lisp
WITH
    workers = 4,
    concurrency = 2
```

Meaning:

- `workers` → total worker threads
- `concurrency` → how many tables at the same time

8️⃣ What happens if threads are too many?

❌ Too many threads causes:

- PostgreSQL connection exhaustion
- Disk IO saturation
- Slower migration (counter-productive)

✔ Best practice:

```nginx
threads ≈ CPU cores
```

9️⃣ Failure handling with threads

If one thread fails:

- Other threads continue
- Error is logged
- Migration does NOT crash completely (unless fatal)

This is called **partial fault tolerance**.

🔟 Mental model (remember forever)

> **pgloader threads are independent workers that read, transform, and COPY data in parallel, while a central controller coordinates constraints and order.**

If you understand:
✔ Thread pool
✔ One thread = one COPY stream
✔ Parallel tables / chunks
✔ Constraints disabled globally

👉 You have **complete mastery of pgloader threading**.

1️⃣1️⃣ Short real-world analogy

Imagine:

- 1 warehouse (PostgreSQL)
- 5 trucks (threads)
- Each truck brings goods (data)
- Security gates (constraints) are opened **after all trucks arrive**

That's exactly pgloader threading.

If you want next, I can:

- Explain **workers vs concurrency** with diagrams
- Explain **pgloader performance tuning**
- Explain **thread behavior during FK-heavy schemas**
- Compare **pgloader threads vs Java threads**

Just tell me 👍